

Etcd Cluster Setup

Configuration of etcd using StateFulset:

```
apiVersion: v1
kind: Service
metadata:
  name: etcd
spec:
  type: ClusterIP
  ports:
    - port: 2379
      name: client
    - port: 2380
      name: peer
  selector:
    app: etcd
---
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: etcd
  labels:
    app: etcd
spec:
  serviceName: etcd
  replicas: 3
  selector:
    matchLabels:
      app: etcd
  template:
    metadata:
      labels:
        app: etcd
    spec:
      containers:
        - name: etcd
          image: quay.io/coreos/etcd:latest
```

```

ports:
- containerPort: 2379
  name: client
- containerPort: 2380
  name: peer
volumeMounts:
- name: data
  mountPath: /var/run/etcd
command:
- /bin/sh
- -c
- |
  PEERS="etcd-0=http://etcd-0.etcd:2380,etcd-1=http://etcd-
1.etcd:2380,etcd-2=http://etcd-2.etcd:2380"
  exec etcd --name ${HOSTNAME} \
    --listen-peer-urls http://0.0.0.0:2380 \
    --listen-client-urls http://0.0.0.0:2379 \
    --advertise-client-urls http://${HOSTNAME}.etcd:2379 \
    --initial-advertise-peer-urls http://${HOSTNAME}:2380 \
    --initial-cluster-token etcd-cluster-1 \
    --initial-cluster ${PEERS} \
    --initial-cluster-state new \
    --data-dir /var/run/etcd/default.etcd
volumeClaimTemplates:
- metadata:
    name: data
  spec:
    storageClassName: longhorn
    accessModes: [ "ReadWriteOnce" ]
    resources:
      requests:
        storage: 1Gi

```

In above configuration we are creating a service for internal communication of etcd pods (happens on port 2379) and external communication between patroni and etcd (happens on port 2380). A statefulset of 3 replicas is configured, with configuration for etcd. Persistent volume (PV) of 1GB for each pod is claimed using PVC.

```
kubectl apply -f file_name.yaml
```

Patroni Cluster Setup

```
apiVersion: v1
kind: Service
metadata:
  name: patroni
  namespace: default
spec:
  type: LoadBalancer
  ports:
    - name: postgresql
      port: 5432
  selector:
    app: patroni
```

Service of load balancer type is created for patroni cluster and external IP is assigned from IpAddress-Pool. Below configurations are for IpAddress-Pool.

```
apiVersion: v1
kind: ConfigMap
metadata:
  namespace: metallb-system
  name: config
data:
  config: |
    address-pools:
    - name: default
      protocol: layer2
      addresses:
      - 10.1.41.206-10.1.41.210
  ---
```

```
apiVersion: metallb.io/v1beta1
kind: L2Advertisement
metadata:
  name: empty
  namespace: metallb-system
```

IpAddress-pool assign the ip from addresses range. When a service in Kubernetes is assigned an external IP, MetalLB running in L2 mode advertises that IP on the network so that devices in the same Layer 2 domain (like switches and routers) know how to reach it.
(Without L2Advertisement we can not access the patroni cluster outside Kubernetes cluster).

Configuration for patroni (ConfigMap)

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: patroni-config
  namespace: default
data:
  patroni.yml: |
    scope: postgres-ha
    namespace: /service/
    name: patroni

    etcd:
      hosts: etcd-0.etcd:2379,etcd-1.etcd:2379,etcd-2.etcd:2379

    bootstrap:
      dcs:
        ttl: 30
        loop_wait: 10
        retry_timeout: 10
        maximum_lag_on_failover: 1048576
        postgresql:
          use_pg_rewind: true
          parameters:
            max_connections: 100
            max_locks_per_transaction: 64
            max_worker_processes: 8
            wal_level: replica
            hot_standby: "on"
            wal_keep_size: 1024
            archive_mode: "on"
            archive_timeout: 1800s

    postgresql:
```

```
listen: "*"
connect_address:
"${HOSTNAME}.patroni.${PATRONI_KUBERNETES_NAMESPACE}.svc.cluster.local:5432"
data_dir: /var/lib/postgresql/data
bin_dir: /usr/lib/postgresql/15/bin
authentication:
  superuser:
    username: postgres
    password: postgres
  replication:
    username: replicator
    password: replicator
parameters:
  archive_mode: "on"
  archive_command: 'cp %p /var/lib/postgresql/archive/%f'

tags:
  nofailover: false
  noloadbalance: false
  clonefrom: false
```

Config map is used to store essential configuration for patroni pods.

Statefulset for Patroni

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: patroni
  namespace: default
spec:
  serviceName: "patroni"
  replicas: 2
  selector:
    matchLabels:
      app: patroni
  template:
    metadata:
      labels:
```

```

    app: patroni
spec:
  containers:
    - name: patroni
      image: adeee11/patroni:latest
      ports:
        - containerPort: 5432
      env:
        - name: PATRONI_KUBERNETES_NAMESPACE
          valueFrom:
            fieldRef:
              fieldPath: metadata.namespace
        - name: PATRONI_KUBERNETES_USE_ENDPOINTS
          value: "true"
        - name: PATRONI_ETCD_HOSTS
          value: "etcd-0.etcd:2379,etcd-1.etcd:2379,etcd-2.etcd:2379"
      volumeMounts:
        - name: postgresql-data
          mountPath: /var/lib/postgresql/data
        - name: config
          mountPath: /etc/patroni
          subPath: patroni.yml
  volumes:
    - name: config
      configMap:
        name: patroni-config
volumeClaimTemplates:
  - metadata:
      name: postgresql-data
    spec:
      storageClassName: longhorn
      accessModes: [ "ReadWriteOnce" ]
      resources:
        requests:
          storage: 1Gi

```

A statefulset of 2 replicas is configured, attached the configMap using volume for patroni and configuration. Persistent volume (PV) of 1GB for each pod is claimed using PVC.

```
kubectl apply -f <name>.yaml
```

Accessing the Cluster

We can access the cluster using external ip (Load Balancer) assigned by IpAddress-pool using following command.

```
psql -h <external-ip> -p 5432 -U postgres
```

```
Password = postgres
```

If Load Balancer connect you to replicator (Read-only pod) for first time, exit and login back you. For second time you will be connected to leader (Read-write pod). This is how Load is balanced.

That All!